

Day29_EDA_WITH_LLM

June 27, 2025

1 LLM-Powered Exploratory Data Analysis (EDA) using Python, Pandas, Gradio & Mistral

1.1 Objective

In this notebook, we will: - Perform manual EDA using Python and seaborn - Identify missing values and distributions - Use a Large Language Model (LLM) to generate intelligent EDA insights - Build a Gradio-based AI-powered web app for EDA automation

1.2 Install Required Libraries

We will use: - pandas for data handling - matplotlib and seaborn for visualizations - gradio for the web interface - ollama for AI-generated insights using Mistral LLM

```
[2]: # Install requiried packages if dont have
! pip install pandas seaborn matplotlib gradio ollama
```

```
Requirement already satisfied: pandas in c:\users\lenovo\anaconda3\lib\site-packages (2.2.3)
```

```
Requirement already satisfied: seaborn in c:\users\lenovo\anaconda3\lib\site-packages (0.13.2)
```

```
Requirement already satisfied: matplotlib in c:\users\lenovo\anaconda3\lib\site-packages (3.10.0)
```

```
Requirement already satisfied: gradio in c:\users\lenovo\anaconda3\lib\site-packages (5.34.2)
```

```
Requirement already satisfied: ollama in c:\users\lenovo\anaconda3\lib\site-packages (0.5.1)
```

```
Requirement already satisfied: numpy>=1.26.0 in c:\users\lenovo\anaconda3\lib\site-packages (from pandas) (2.1.3)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\lenovo\anaconda3\lib\site-packages (from pandas) (2.9.0.post0)
```

```
Requirement already satisfied: pytz>=2020.1 in c:\users\lenovo\anaconda3\lib\site-packages (from pandas) (2024.1)
```

```
Requirement already satisfied: tzdata>=2022.7 in c:\users\lenovo\anaconda3\lib\site-packages (from pandas) (2025.2)
```

```
Requirement already satisfied: contourpy>=1.0.1 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (1.3.1)
```

Requirement already satisfied: cyclер>=0.10 in
c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (4.55.3)

Requirement already satisfied: kiwisolver>=1.3.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (1.4.8)

Requirement already satisfied: packaging>=20.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (24.2)

Requirement already satisfied: pillow>=8 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (11.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib) (3.2.0)

Requirement already satisfied: aiofiles<25.0,>=22.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (24.1.0)

Requirement already satisfied: anyio<5.0,>=3.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (4.7.0)

Requirement already satisfied: audioop-lts<1.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.2.1)

Requirement already satisfied: fastapi<1.0,>=0.115.2 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.115.13)

Requirement already satisfied: ffmpeg in c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.6.0)

Requirement already satisfied: gradio-client==1.10.3 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (1.10.3)

Requirement already satisfied: groovy~=0.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.1.2)

Requirement already satisfied: httpx>=0.24.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.28.1)

Requirement already satisfied: huggingface-hub>=0.28.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.33.1)

Requirement already satisfied: jinja2<4.0 in c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (3.1.6)

Requirement already satisfied: markupsafe<4.0,>=2.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (3.0.2)

Requirement already satisfied: orjson~=3.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (3.10.18)

Requirement already satisfied: pydantic<2.12,>=2.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (2.10.3)

Requirement already satisfied: pydub in c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.25.1)

Requirement already satisfied: python-multipart>=0.0.18 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.0.20)

Requirement already satisfied: pyyaml<7.0,>=5.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (6.0.2)

Requirement already satisfied: ruff>=0.9.3 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.12.0)

Requirement already satisfied: safehttpx<0.2.0,>=0.1.6 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.1.6)

Requirement already satisfied: semantic-version~=2.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (2.10.0)

Requirement already satisfied: starlette<1.0,>=0.40.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.46.2)

Requirement already satisfied: tomlkit<0.14.0,>=0.12.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.13.2)

Requirement already satisfied: typer<1.0,>=0.12 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.16.0)

Requirement already satisfied: typing-extensions~=4.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (4.12.2)

Requirement already satisfied: uvicorn>=0.14.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio) (0.34.3)

Requirement already satisfied: fsspec in c:\users\lenovo\anaconda3\lib\site-
packages (from gradio-client==1.10.3->gradio) (2025.3.2)

Requirement already satisfied: websockets<16.0,>=10.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from gradio-client==1.10.3->gradio)
(15.0.1)

Requirement already satisfied: idna>=2.8 in c:\users\lenovo\anaconda3\lib\site-
packages (from anyio<5.0,>=3.0->gradio) (3.7)

Requirement already satisfied: sniffio>=1.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from anyio<5.0,>=3.0->gradio)
(1.3.0)

Requirement already satisfied: annotated-types>=0.6.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from pydantic<2.12,>=2.0->gradio)
(0.6.0)

Requirement already satisfied: pydantic-core==2.27.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from pydantic<2.12,>=2.0->gradio)
(2.27.1)

Requirement already satisfied: click>=8.0.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from typer<1.0,>=0.12->gradio)
(8.1.8)

Requirement already satisfied: shellingham>=1.3.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from typer<1.0,>=0.12->gradio)
(1.5.0)

Requirement already satisfied: rich>=10.11.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from typer<1.0,>=0.12->gradio)
(13.9.4)

Requirement already satisfied: colorama in c:\users\lenovo\anaconda3\lib\site-
packages (from click>=8.0.0->typer<1.0,>=0.12->gradio) (0.4.6)

Requirement already satisfied: certifi in c:\users\lenovo\anaconda3\lib\site-
packages (from httpx>=0.24.1->gradio) (2025.4.26)

Requirement already satisfied: httpcore==1.* in
c:\users\lenovo\anaconda3\lib\site-packages (from httpx>=0.24.1->gradio) (1.0.9)

Requirement already satisfied: h11>=0.16 in c:\users\lenovo\anaconda3\lib\site-
packages (from httpcore==1.*->httpx>=0.24.1->gradio) (0.16.0)

Requirement already satisfied: filelock in c:\users\lenovo\anaconda3\lib\site-
packages (from huggingface-hub>=0.28.1->gradio) (3.17.0)

Requirement already satisfied: requests in c:\users\lenovo\anaconda3\lib\site-

```

packages (from huggingface-hub>=0.28.1->gradio) (2.32.3)
Requirement already satisfied: tqdm>=4.42.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from huggingface-
hub>=0.28.1->gradio) (4.67.1)
Requirement already satisfied: six>=1.5 in c:\users\lenovo\anaconda3\lib\site-
packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from
rich>=10.11.0->typer<1.0,>=0.12->gradio) (2.2.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
c:\users\lenovo\anaconda3\lib\site-packages (from
rich>=10.11.0->typer<1.0,>=0.12->gradio) (2.19.1)
Requirement already satisfied: mdurl~=0.1 in c:\users\lenovo\anaconda3\lib\site-
packages (from markdown-it-py>=2.2.0->rich>=10.11.0->typer<1.0,>=0.12->gradio)
(0.1.0)
Requirement already satisfied: charset-normalizer<4,>=2 in
c:\users\lenovo\anaconda3\lib\site-packages (from requests->huggingface-
hub>=0.28.1->gradio) (3.3.2)
Requirement already satisfied: urllib3<3,>=1.21.1 in
c:\users\lenovo\anaconda3\lib\site-packages (from requests->huggingface-
hub>=0.28.1->gradio) (2.3.0)

```

1.3 Load the Titanic Dataset

We'll start by loading the dataset into a pandas DataFrame.

```

[3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load Titanic Dataset
url = r"E:\Data backup 25.6.25\Downloads\titanic_ dataset_final.csv"
df = pd.read_csv(url)
df

```

```

[3]:
   PassengerId  Survived  Pclass \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
..          ...     ...     ...
886          887         0       2
887          888         1       1
888          889         0       3
889          890         1       1
890          891         0       3

```

		Name	Sex	Age	SibSp	\
0		Braund, Mr. Owen Harris	male	22.0	1	
1	Cummings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1	
4		Allen, Mr. William Henry	male	35.0	0	
..			
886		Montvila, Rev. Juozas	male	27.0	0	
887		Graham, Miss. Margaret Edith	female	19.0	0	
888	Johnston, Miss. Catherine Helen "Carrie"		female	NaN	1	
889		Behr, Mr. Karl Howell	male	26.0	0	
890		Dooley, Mr. Patrick	male	32.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..	
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

1.4 View Statistical Summary

This gives basic descriptive statistics like count, mean, std deviation, etc.

```
[4]: print(df.describe())
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429

min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

1.5 Check for Missing Values

This shows the number of missing entries for each column.


```
[5]: print("\nMissing Values:\n", df.isnull().sum())
```

```
Missing Values:
 PassengerId      0
 Survived         0
 Pclass          0
 Name            0
 Sex             0
 Age            177
 SibSp           0
 Parch           0
 Ticket          0
 Fare           0
 Cabin          687
 Embarked        2
dtype: int64
```

1.6 Plot Survival Distribution

We visualize how many passengers survived vs. not survived.

```
[6]: sns.countplot(x='Survived', data=df)
plt.title("Survival Count")
plt.show()
```



1.7 What if AI Could Do All This Automatically?

Let's now use a large language model (Mistral) through `ollama` to generate insights automatically.

1.8 AI-Generated Insights Using Mistral

We pass the summary to the LLM and get automatic interpretation.

```
[8]: import ollama

def generate_insights(df_summary):
    prompt = f"Analyze the dataset summary and provide insights:
    ↪\n\n{df_summary}"
    response = ollama.chat(model="mistral", messages=[{"role": "user",
    ↪"content": prompt}])
    return response['message']['content']

# Generate AI Insights
summary = df.describe().to_string()
insights = generate_insights(summary)
```

```
print("\n -- AI-Generated Insights:\n", insights)
```

-- AI-Generated Insights:

This dataset appears to be the Titanic passenger manifest data, which includes variables such as PassengerID, Survived (1 if survived, 0 if not), Pclass (ticket class), Age, SibSp (number of siblings or spouses aboard the ship), Parch (number of parents or children aboard), and Fare.

Here are some insights from the dataset summary:

1. The total number of passengers is 891. There were no missing values for any variable in this dataset.
2. The mean age of passengers is approximately 29.7 years, with a standard deviation of around 14.53 years. This suggests a wide distribution of ages among the passengers.
3. Majority of the passengers survived (mean Survived = 0.38), but there were still many who did not survive (mean Survived = 0.62 would mean all passengers survived).
4. Most passengers traveled in 3rd class (mean Pclass = 2.31, median Pclass = 3.0), suggesting a high number of passengers traveling in the cheapest class.
5. On average, each passenger had about 0.5 family members or travel companions aboard the ship (SibSp + Parch).
6. The fare varies greatly, with a mean of around 32.20 and a standard deviation of approximately 49.69. This indicates that there was significant disparity in ticket prices among passengers.
7. The youngest passenger was only 0.42 years old, while the oldest was 80 years old. The median age is 28, with 25% and 75% of passengers being younger (20.125) and older (38) than this value, respectively.
8. Interestingly, despite the large number of passengers in 3rd class, it seems that a non-zero percentage of first-class passengers did not survive (mean Survived = 0.68 for Pclass=1). This could potentially be an interesting factor to explore further when analyzing survival rates based on different ticket classes.

1.9 Create a Basic Web App for EDA Using Gradio

This app lets users upload a CSV file and get insights.

```
[9]: import gradio as gr
import pandas as pd
```



```

def eda_analysis(file):
    df = pd.read_csv(file.name)
    summary = df.describe().to_string()
    insights = generate_insights(summary) # This should be defined separately
    return insights

# Create Web Interface
demo = gr.Interface(
    fn=eda_analysis,
    inputs=gr.File(label="Upload your CSV file", file_types=[".csv"]),
    outputs=gr.Textbox(label="AI-Generated Insights", lines=20),
    title="AI-Powered EDA with Mistral",
    description="""
    <div style="text-align: center;">
        <h3 style="color: #4CAF50; margin-bottom: 0;">Explore Your Dataset
    ↪Instantly</h3>
        <p style="color: #555;">Upload a CSV file and get an instant,
    ↪intelligent summary and analysis powered by AI.</p>
    </div>
    """,
    article="""
    <p style="text-align: center; font-size: 14px; color: gray;">
        Built using <strong>Pandas</strong> and <strong>Gradio</strong>. AI
    ↪insights are generated by a custom model powered by Mistral.
        <br>Developed for automated exploratory data analysis (EDA).
    </p>
    """,
    theme="huggingface" # You can also try "default", "soft", or "grass"
)

# Launch App
demo.launch(share=True)

```

C:\Users\Lenovo\anaconda3\Lib\site-packages\gradio\blocks.py:1180: UserWarning: Cannot load huggingface. Caught Exception: 404 Client Error: Not Found for url: https://huggingface.co/api/spaces/huggingface (Request ID: Root=1-685e9283-3fb0de2e00eacb090c3e543f;d93e573a-7699-4138-933b-dea9a97fbd20)

Sorry, we can't find the page you are looking for.

warnings.warn(f"Cannot load {theme}. Caught Exception: {str(e)}")

* Running on local URL: http://127.0.0.1:7860

* Running on public URL: https://34a9e0105dc7e46b81.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging Face Spaces (https://huggingface.co/spaces)

<IPython.core.display.HTML object>

[9]:

1.10 Final Web App with Summary, Missing Values, Insights, and Visuals

This version also adds histograms and a correlation heatmap automatically.

```
[12]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import gradio as gr
import os

# Generate Insights (placeholder)
def generate_insights(summary):
    return "Key Insights:\n- Dataset has been summarized.\n- Add_\n- model-generated insights here."

# Create visualizations
def generate_visualizations(df):
    plot_paths = []
    os.makedirs("plots", exist_ok=True)

    for col in df.select_dtypes(include=['number']).columns:
        plt.figure(figsize=(6, 4))
        sns.histplot(df[col], bins=20, kde=True, color="skyblue")
        plt.title(f"Distribution of {col}")
        plot_path = f"plots/{col}_hist.png"
        plt.savefig(plot_path)
        plot_paths.append(plot_path)
        plt.close()

    return plot_paths

# EDA function
def eda_analysis(file):
    df = pd.read_csv(file.name)

    # Fill missing values
    for col in df.select_dtypes(include=['number']).columns:
        df[col].fillna(df[col].median(), inplace=True)

    for col in df.select_dtypes(include=['object']).columns:
        df[col].fillna(df[col].mode()[0], inplace=True)

    # Summary
    summary = df.describe(include='all').to_string()
```

```

missing = df.isnull().sum().to_string()

# Insights and visualizations
insights = generate_insights(summary)
plot_paths = generate_visualizations(df)

report = f"Summary Statistics:\n\n{summary}\n\nMissing Values:
↪\n{missing}\n\nAI Insights:\n{insights}"

return report, plot_paths

# Gradio Interface
demo = gr.Interface(
    fn=eda_analysis,
    inputs=gr.File(label="Upload your CSV file", file_types=[".csv"]),
    outputs=[
        gr.Textbox(label="EDA Report", lines=30),
        gr.Gallery(label="Visualizations")
    ],
    title="AI-Powered Exploratory Data Analysis",
    description="""
    Upload a CSV file to get automatic EDA insights. This includes summary_
    ↪statistics, missing value handling, and visualizations.
    """,
    article="""
    Built using Pandas, Seaborn, and Gradio. You can enhance this with custom_
    ↪models for deeper insights.
    """,
    theme="huggingface"
)

demo.launch(share=True)

```

C:\Users\Lenovo\anaconda3\Lib\site-packages\gradio\blocks.py:1180: UserWarning: Cannot load huggingface. Caught Exception: 404 Client Error: Not Found for url: https://huggingface.co/api/spaces/huggingface (Request ID: Root=1-685e9741-58cddf7b65c7c13251909fb9;24b70883-cefd-4fbd-be16-5a8fc21deaec)

Sorry, we can't find the page you are looking for.

warnings.warn(f"Cannot load {theme}. Caught Exception: {str(e)}")

* Running on local URL: http://127.0.0.1:7862

* Running on public URL: https://0d85562b21700e438a.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging Face Spaces (<https://huggingface.co/spaces>)

<IPython.core.display.HTML object>

[12]:

```
C:\Users\Lenovo\AppData\Local\Temp\ipykernel_8108\1815311375.py:33:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].median(), inplace=True)
C:\Users\Lenovo\AppData\Local\Temp\ipykernel_8108\1815311375.py:36:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

1.11 Conclusion

With just a few lines of code, we've: - Explored the Titanic dataset manually - Generated AI-powered summaries and visualizations - Built two different web apps using Gradio: one simple, one advanced - Integrated LLMs like Mistral using Ollama for smart insights

This is a scalable, interactive approach to rapid data exploration in real-world projects.

[]: