

Day42_SRL_House_Price_Prediction_Project

July 11, 2025

Simple Linear Regression: House Price Prediction Project

Project Description

This project uses **Simple Linear Regression** to predict house prices based on the **square footage of living area (sqft_living)** using a real dataset.

We'll go through:

- Data loading and exploration
- Feature selection
- Visualizing the relationship
- Training a Linear Regression model
- Evaluating model performance
- Making predictions on unseen data

Objective

Build a regression model to understand the relationship between square footage and price, and use it to predict prices for new, unseen homes.

1 Import Required Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
```

2 Load the Dataset

We are loading the dataset from the given path using `pandas.read_csv()` and displaying the first few rows using `df.head()`.

```
[2]: df = pd.read_csv(r"D:\Assignment Practice\M3\08 July\08 July\7th- slr\SLR - 1
↳House price prediction\House_data.csv")
df.head()
```

```
[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900.0	3	1.00	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	
2	5631500400	20150225T000000	180000.0	2	1.00	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	\
0	5650	1.0	0	0	...	7	1180	0	
1	7242	2.0	0	0	...	7	2170	400	
2	10000	1.0	0	0	...	6	770	0	
3	5000	1.0	0	0	...	7	1050	910	
4	8080	1.0	0	0	...	8	1680	0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	

	sqft_lot15
0	5650
1	7639
2	8062
3	5000
4	7503

[5 rows x 21 columns]

3 Understand the Data

info() gives an overview of column types and non-null counts.

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              21613 non-null  int64
1   date            21613 non-null  object
2   price           21613 non-null  float64
3   bedrooms        21613 non-null  int64
4   bathrooms       21613 non-null  float64
5   sqft_living     21613 non-null  int64
6   sqft_lot        21613 non-null  int64
```

```

7  floors          21613 non-null float64
8  waterfront      21613 non-null int64
9  view            21613 non-null int64
10 condition       21613 non-null int64
11 grade           21613 non-null int64
12 sqft_above      21613 non-null int64
13 sqft_basement   21613 non-null int64
14 yr_built        21613 non-null int64
15 yr_renovated    21613 non-null int64
16 zipcode         21613 non-null int64
17 lat             21613 non-null float64
18 long            21613 non-null float64
19 sqft_living15   21613 non-null int64
20 sqft_lot15      21613 non-null int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB

```

describe() gives statistical summary of numerical features.

```
[4]: df.describe()
```

```

[4]:
count      id      price      bedrooms      bathrooms      sqft_living  \
count  2.161300e+04  2.161300e+04  21613.000000  21613.000000  21613.000000
mean    4.580302e+09  5.401822e+05    3.370842    2.114757    2079.899736
std     2.876566e+09  3.673622e+05    0.930062    0.770163    918.440897
min     1.000102e+06  7.500000e+04    0.000000    0.000000    290.000000
25%     2.123049e+09  3.219500e+05    3.000000    1.750000    1427.000000
50%     3.904930e+09  4.500000e+05    3.000000    2.250000    1910.000000
75%     7.308900e+09  6.450000e+05    4.000000    2.500000    2550.000000
max     9.900000e+09  7.700000e+06    33.000000    8.000000   13540.000000

count      sqft_lot      floors      waterfront      view      condition  \
count  2.161300e+04  21613.000000  21613.000000  21613.000000  21613.000000
mean    1.510697e+04    1.494309    0.007542    0.234303    3.409430
std     4.142051e+04    0.539989    0.086517    0.766318    0.650743
min     5.200000e+02    1.000000    0.000000    0.000000    1.000000
25%     5.040000e+03    1.000000    0.000000    0.000000    3.000000
50%     7.618000e+03    1.500000    0.000000    0.000000    3.000000
75%     1.068800e+04    2.000000    0.000000    0.000000    4.000000
max     1.651359e+06    3.500000    1.000000    4.000000    5.000000

count      grade      sqft_above      sqft_basement      yr_built      yr_renovated  \
count  21613.000000  21613.000000  21613.000000  21613.000000  21613.000000
mean         7.656873   1788.390691    291.509045   1971.005136    84.402258
std         1.175459    828.090978   442.575043    29.373411   401.679240
min          1.000000    290.000000    0.000000   1900.000000    0.000000
25%          7.000000   1190.000000    0.000000   1951.000000    0.000000
50%          7.000000   1560.000000    0.000000   1975.000000    0.000000

```

75%	8.000000	2210.000000	560.000000	1997.000000	0.000000
max	13.000000	9410.000000	4820.000000	2015.000000	2015.000000

	zipcode	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	98077.939805	47.560053	-122.213896	1986.552492	12768.455652
std	53.505026	0.138564	0.140828	685.391304	27304.179631
min	98001.000000	47.155900	-122.519000	399.000000	651.000000
25%	98033.000000	47.471000	-122.328000	1490.000000	5100.000000
50%	98065.000000	47.571800	-122.230000	1840.000000	7620.000000
75%	98118.000000	47.678000	-122.125000	2360.000000	10083.000000
max	98199.000000	47.777600	-121.315000	6210.000000	871200.000000

4 Select Features for Simple Linear Regression

We're using only one feature `sqft_living` to predict the house price, hence it's a Simple Linear Regression.


```
[5]: X = np.array(df['sqft_living']).reshape(-1, 1) # Independent variable
     y = np.array(df['price'])                    # Dependent variable
```

- sklearn models expect X to be a 2D array → shape (n_samples, n_features)
- `.reshape(-1, 1)` converts a 1D array to 2D
- `-1` lets NumPy figure out the number of rows
- `1` means one feature (column)
- `y` can stay 1D — it only contains target values.

5 Visualize the Relationship

A scatter plot helps to visualize whether there's a linear relationship between `sqft_living` and `price`.

```
[6]: plt.scatter(X, y, color='blue', alpha=0.5)
     plt.title("House Price vs. Square Feet Living")
     plt.xlabel("Square Feet Living")
     plt.ylabel("Price")
     plt.show()
```



6 Split the Data

We split the data into training and testing sets (80-20 split) so that we can evaluate the model's performance on unseen data.

```
[7]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

7 Build and Train the Model

We create a Linear Regression model and fit it using the training data.

```
[8]: model = LinearRegression()  
     model.fit(X_train, y_train)
```

```
[8]: LinearRegression()
```

8 Get Model Parameters

- **Intercept ()**: The constant term — it's the value of the price when `sqft_living = 0`.
- **Coefficient ()**: The slope — it tells how much the price increases for each additional square foot.

Equation of the model:

$$\text{Predicted Price} = b_0 + b_1 \times \text{sqft_living}$$

```
[9]: print("Intercept (b):", model.intercept_)
      print("Coefficient (b):", model.coef_[0])
```

```
Intercept (b): -42291.12839802564
Coefficient (b): 279.7397784623452
```

9 Predict and Evaluate the Model

- MSE measures average squared difference between actual and predicted values.
- R^2 Score tells how well the model explains the variation in the target variable.

```
[10]: y_pred = model.predict(X_test)
      print(y_pred)
```

```
[ 536770.21301903  768954.22914278 1012327.83640502 ...  638595.49237932
 587123.37314225  676640.1022502 ]
```

9.1 Metrics


```
[11]: mse = mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)
```

```
[12]: print("Mean Squared Error:", mse)
      print("R2 Score:", r2)
```

```
Mean Squared Error: 76570251342.04109
R2 Score: 0.4941006145983624
```

10 Plot the Regression Line

```
[13]: plt.scatter(X_test, y_test, color='blue', alpha=0.5, label="Actual")
      plt.plot(X_test, y_pred, color='red', linewidth=2, label="Regression Line")
      plt.title("Regression Line - House Price Prediction")
      plt.xlabel("Square Feet Living")
      plt.ylabel("Price")
      plt.legend()
      plt.show()
```



11 Predict on New (Unseen) Data

We use the trained model to predict house prices for unseen data points.

```
[14]: # Example: Predict price for 2000, 2500, and 3000 square feet
new_sqft = np.array([[2000], [2500], [3000]])
predicted_price = model.predict(new_sqft)

for sqft, price in zip(new_sqft, predicted_price):
    print(f"Predicted price for {sqft[0]} sqft = ${price:,.2f}")
```

Predicted price for 2000 sqft = \$517,188.43

Predicted price for 2500 sqft = \$657,058.32

Predicted price for 3000 sqft = \$796,928.21

12 Conclusion

- We built a Simple Linear Regression model using sqft_living to predict house prices.
- Evaluated it using metrics like MSE and R^2 .

- Plotted the regression line and predicted new values.

[]: