

Day74_Deep_Learning_VanishingGradient_Dropout_Optimization_LossFun

August 28, 2025

1 Deep Learning (Vanishing Gradient, Dropout, Optimization, Loss Functions)

Welcome back to my Deep Learning documentation!

In this notebook, we continue from **Day73**, focusing on important training concepts in neural networks:

- Vanishing Gradient Problem
- Chain Rule in Backpropagation
- Dropout Neurons
- Optimization Techniques
- Loss Functions

2 Vanishing Gradient Problem

Deep neural networks often face the **Vanishing Gradient Problem** during training.

2.1 What is it?

- In deep networks, gradients become **very small** as they propagate backward.
- This makes weight updates negligible in early layers → **network stops learning**.

2.2 Why does it happen?

- Sigmoid and Tanh activations squash values into a small range.
- Their derivatives are small:

$$\sigma'(x) = \sigma(x)(1 - \sigma(x)) \in (0, 0.25)$$

$$\tanh'(x) = 1 - \tanh^2(x) \in (0, 1)$$

- Multiplying many small derivatives \rightarrow gradient approaches 0.

2.3 Symptoms

- Accuracy stops improving after a few epochs.
- Model seems “stuck” at some accuracy.

2.4 Solutions

1. Use **ReLU** / **Leaky ReLU** instead of Sigmoid/Tanh.
2. Apply **Batch Normalization** \rightarrow keeps activations stable.
3. Use better **optimizers** (Adam, RMSProp).
4. **Dropout Neurons** \rightarrow improves generalization.

3 Chain Rule in Backpropagation

Backpropagation is based on the **Chain Rule of Calculus**.

3.1 Formula

If:

$$y = f(g(x))$$

Then derivative:

$$\frac{dy}{dx} = f'(g(x)) \cdot g'(x)$$

3.2 In Neural Networks

- Error flows backward from **Output** \rightarrow **Hidden** \rightarrow **Input**.
- Each layer’s gradient is computed as a product of partial derivatives.

Example (simple 2-layer NN):

$$L = f(z), \quad z = g(h), \quad h = w \cdot x$$

$$\frac{dL}{dw} = \frac{dL}{dz} \cdot \frac{dz}{dh} \cdot \frac{dh}{dw}$$

This shows how gradients are **chained** across layers.

4 Dropout Neurons

Dropout is a **regularization technique** used to prevent overfitting.

4.1 What is Dropout?

- During training, **random neurons are dropped** (set to 0).
- This forces the network to not depend on specific neurons.

4.2 Example

- Suppose a layer has 100 neurons.
- If dropout = 0.5 \rightarrow only ~50 neurons are active at each step.

4.3 At Inference (Testing)

- No dropout is applied.
- Weights are **scaled** to match training conditions.

Dropout improves generalization and reduces overfitting.

5 Optimization in Deep Learning

Optimizers control **how weights are updated** during training.

5.1 Types of Gradient Descent

5.1.1 Stochastic Gradient Descent (SGD)

- Updates weights after **each sample**.
- Faster but noisy.

5.1.2 Batch Gradient Descent (BGD)

- Uses the **entire dataset** for each update.
- Very accurate but slow.

5.1.3 Mini-Batch Gradient Descent

- Uses small batches of data.
- Default choice in practice.

5.2 Advanced Optimizers

5.2.1 Adam (Adaptive Moment Estimation)

- Combines **momentum** + **adaptive learning rate**.
- Very popular for deep learning.

5.2.2 Adamax

- Variant of Adam.
- Works better with **sparse gradients**.

5.2.3 Adadelta

- Dynamically adjusts learning rate.

5.2.4 RMSProp

- Keeps moving average of squared gradients.
- Prevents exploding/vanishing updates.

6 Loss Functions

Loss functions measure **how far predictions are from actual values**.

6.1 Regression Losses

6.1.1 Mean Absolute Error (MAE)

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_{true}^{(i)} - y_{pred}^{(i)}|$$

6.1.2 Mean Squared Error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_{true}^{(i)} - y_{pred}^{(i)})^2$$

6.1.3 Log Loss

- Used for probabilistic regression.

6.2 Classification Losses

6.2.1 Binary Cross-Entropy

For binary classification (0/1):

$$Loss = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i) \right]$$

6.2.2 Categorical Cross-Entropy

For multi-class classification (one-hot labels):

$$Loss = -\sum_{i=1}^n \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$$

6.2.3 Sparse Categorical Cross-Entropy

- Similar to categorical cross-entropy, but labels are **integers** instead of one-hot.

7 Key Insights

- Backpropagation = **Forward Propagation + Chain Rule + Weight Updates**.
- If weights stop updating → **Vanishing Gradient Problem**.
- **Dropout Neurons + L1/L2 Regularization** help reduce overfitting.
- Advanced optimizers (Adam, RMSProp, etc.) improve convergence speed.
- Different loss functions are used for **regression vs classification** tasks.

8 Summary

- Vanishing Gradient slows training → solved with ReLU, Dropout, BatchNorm, better optimizers.
- Chain Rule = backbone of backpropagation.
- Dropout Neurons = prevent overfitting.
- Optimizers = decide how learning happens.
- Loss Functions = measure error in regression/classification.

9 Conclusion

In this notebook, we covered:

- Vanishing Gradient Problem
- Chain Rule in Backpropagation
- Dropout Neurons

- Optimization Methods (SGD, Adam, RMSProp, etc.)
- Loss Functions (Regression & Classification)

These concepts help us train **deeper and more accurate neural networks**.