

Day48_Support_Vector_Regression_(SVR) _Salary_Prediction

July 17, 2025

Today, we are learning about the **Support Vector Regression (SVR)** algorithm – a powerful regression technique especially useful when the data doesn't follow a straight line.

We'll implement SVR step-by-step using Python and visualize how different configurations affect the predictions.

What We'll Cover:

1. **Simple Implementation (Beginner-friendly)**
Build SVR models with different kernels **without using functions** to understand the basics.
2. **With Function (Reusable Code)**
Wrap the SVR logic into a function to quickly test multiple configurations (kernel, degree, C, gamma, epsilon).
3. **All-in-One Comparison**
Try **all parameter combinations** (like kernel, degree, C, epsilon, gamma) and build a **summary table** to compare model predictions for Level 6.

Support Vector Regression (SVR)

What is SVR?

- **SVR** stands for **Support Vector Regressor**.
- It is used to **predict continuous values** like salary, price, etc.
- It is built on the same idea as **SVM (Support Vector Machine)**, which is used for classification.
- The main goal of SVR is to **fit a line (hyperplane)** such that it stays within a margin (epsilon) and allows for **some error**.

SVR vs SVM

Concept	SVR (Regression)	SVM (Classification)
Output Type	Continuous (e.g. salary)	Categorical (e.g. yes/no)
Target Variable	Numerical	Binary/Multiclass
Best Fit	Hyperplane	Decision Boundary
Use Case	Predicting values	Separating classes
Goal	Minimize error and stay within margin	Maximize margin between classes

Key Concepts in SVR

- **Hyperplane:** A line or plane that best fits the data.

- **Support Vectors:** Points closest to the hyperplane — these define the margin.
- **Margin:** Distance between the hyperplane and the closest data points.
- **Epsilon ():** Defines a margin of tolerance where no penalty is given for error.
- **Objective:** Keep predictions within \pm margin.

Think of SVR like fitting a straight line — but instead of going through every point, it fits **within a soft margin**, ignoring small errors.

Support Vector Regression (SVR) – Parameters & Tuning Guide

Parameter Descriptions

Parameter Meaning	
kernel	Defines the shape of the model's curve or surface. Options: 'linear', 'poly', 'rbf', 'sigmoid'
degree	Degree of the polynomial kernel (only applies when kernel='poly'). Higher degree = more complex curve
C	Regularization parameter. High C = less error tolerance (stricter margin), low C = more flexibility
epsilon	Acceptable error margin in prediction. No penalty within \pm epsilon. Default is 0.1
gamma	Defines influence of a single training point. Options: 'scale', 'auto', or a float value

Recommended Parameter Values to Try

Parameter	Suggested Values
kernel	'linear', 'poly', 'rbf', 'sigmoid'
degree	2, 3, 4, 5, 6 (<i>only if kernel='poly'</i>)
C	0.1, 1, 10, 100
epsilon	0.01, 0.1, 1
gamma	'scale', 'auto'

Tip: Start with default values and adjust one parameter at a time to see how it impacts performance. For advanced optimization, consider using `GridSearchCV`.

Example: SVR for Salary Prediction

We are predicting salary based on employee level using SVR. The dataset looks like this:

Position	Level	Salary
Jr Software Engineer	1	45000
Sr Software Engineer	2	50000
Team Lead	3	60000
Manager	4	80000
Sr Manager	5	110000
...

1 Simple Implementation (Beginner-friendly)

1.1 Import Required Libraries

```
[24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.svm import SVR
from sklearn.preprocessing import StandardScaler
```

1.2 Load the Dataset

```
[2]: dataset = pd.read_csv(r"C:\Users\Lenovo\Downloads\emp_sal.csv")
dataset.head()
```

```
[2]:
```

	Position	Level	Salary
0	Jr Software Engineer	1	45000
1	Sr Software Engineer	2	50000
2	Team Lead	3	60000
3	Manager	4	80000
4	Sr manager	5	110000

1.3 Define Features and Target

```
[3]: X = dataset.iloc[:, 1:2].values # Level
y = dataset.iloc[:, 2].values # Salary
```

1.4 Feature Scaling (Important for SVR)

Why We Scale X and y?

- **SVR needs scaled data** for better accuracy.
- $X_{\text{scaled}} = \text{sc_X.fit_transform}(X) \rightarrow$ Scales Level (input).
- $y_{\text{scaled}} = \text{sc_y.fit_transform}(y.\text{reshape}(-1, 1)).\text{flatten}()$
 \rightarrow Reshapes y to 2D, scales it, then flattens back to 1D (SVR needs 1D y).

```
[4]: # SVR needs scaling
sc_X = StandardScaler()
sc_y = StandardScaler()
```

```
[5]: X_scaled = sc_X.fit_transform(X) # Scale input feature (Level)
# Reshape y to 2D, scale it, then flatten back to 1D (required by SVR)
y_scaled = sc_y.fit_transform(y.reshape(-1, 1)).flatten() # Scale target
↳ (Salary)
```

1.5 Try Different SVR Models (Without Function , Simple)

1.5.1 SVR with Default (RBF) Kernel

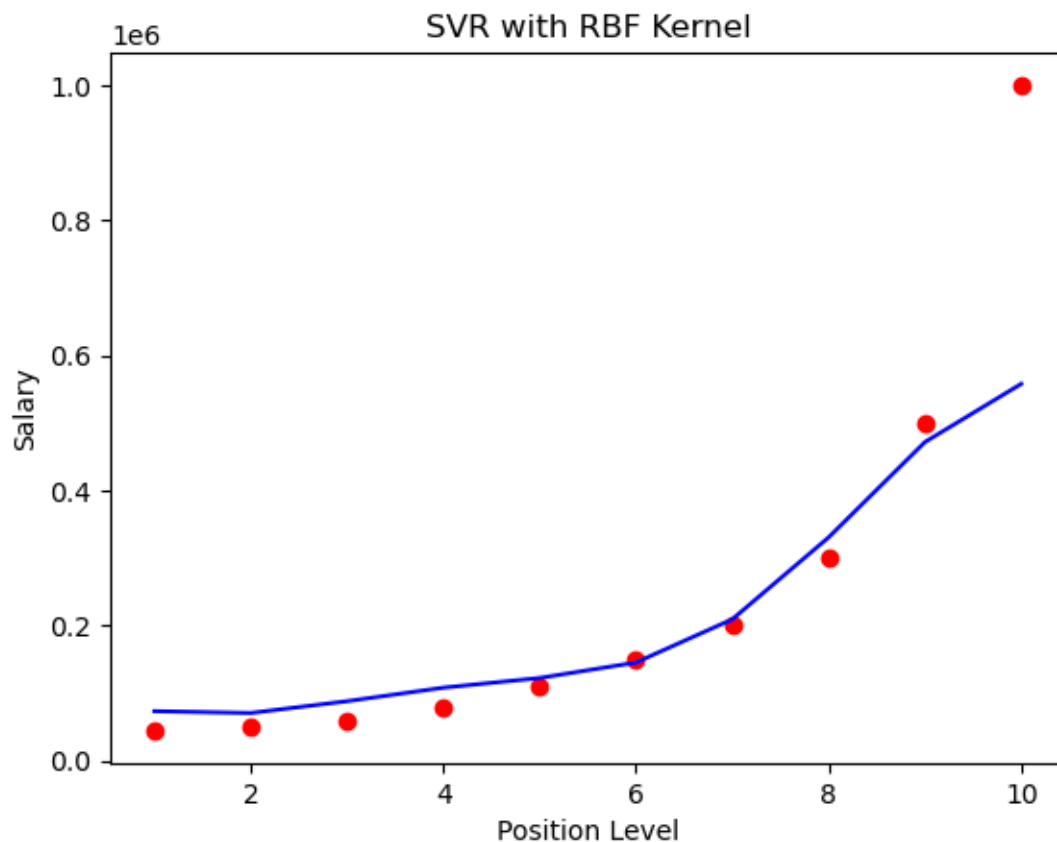
```
[9]: svr_rbf = SVR()  
svr_rbf.fit(X_scaled, y_scaled)
```

```
[9]: SVR()
```

```
[10]: # Predict for level 6  
level_6_scaled = sc_X.transform([[6]])  
pred_rbf_scaled = svr_rbf.predict(level_6_scaled)  
pred_rbf = sc_y.inverse_transform(pred_rbf_scaled.reshape(-1, 1))  
print("SVR (RBF) → Predicted salary for level 6:", round(pred_rbf[0][0], 2))
```

SVR (RBF) → Predicted salary for level 6: 145503.11

```
[11]: plt.scatter(X, y, color='red')  
plt.plot(X, sc_y.inverse_transform(svr_rbf.predict(X_scaled).reshape(-1, 1)),  
        color='blue')  
plt.title("SVR with RBF Kernel")  
plt.xlabel("Position Level")  
plt.ylabel("Salary")  
plt.show()
```



1.5.2 SVR with Polynomial Kernel (Degree 4)

```
[12]: svr_poly4 = SVR(kernel='poly', degree=4)
svr_poly4.fit(X_scaled, y_scaled)

pred_poly4_scaled = svr_poly4.predict(level_6_scaled)
pred_poly4 = sc_y.inverse_transform(pred_poly4_scaled.reshape(-1, 1))
print("SVR (Poly, Degree 4) →", round(pred_poly4[0][0], 2))
```

SVR (Poly, Degree 4) → 130000.0

1.5.3 SVR with Polynomial Kernel (Degree 5)

```
[13]: svr_poly5 = SVR(kernel='poly', degree=5)
svr_poly5.fit(X_scaled, y_scaled)

pred_poly5_scaled = svr_poly5.predict(level_6_scaled)
pred_poly5 = sc_y.inverse_transform(pred_poly5_scaled.reshape(-1, 1))
print("SVR (Poly, Degree 5) →", round(pred_poly5[0][0], 2))
```

SVR (Poly, Degree 5) → 225342.67

1.5.4 SVR with Polynomial Kernel (Degree 6)

```
[14]: svr_poly6 = SVR(kernel='poly', degree=6, gamma='scale')
svr_poly6.fit(X_scaled, y_scaled)

pred_poly6_scaled = svr_poly6.predict(level_6_scaled)
pred_poly6 = sc_y.inverse_transform(pred_poly6_scaled.reshape(-1, 1))
print("SVR (Poly, Degree 6) →", round(pred_poly6[0][0], 2))
```

SVR (Poly, Degree 6) → 130000.0

1.5.5 SVR with Sigmoid Kernel

```
[15]: svr_sigmoid = SVR(kernel='sigmoid', degree=6)
svr_sigmoid.fit(X_scaled, y_scaled)

pred_sigmoid_scaled = svr_sigmoid.predict(level_6_scaled)
pred_sigmoid = sc_y.inverse_transform(pred_sigmoid_scaled.reshape(-1, 1))
print("SVR (Sigmoid) →", round(pred_sigmoid[0][0], 2))
```

SVR (Sigmoid) → 276639.07

1.6 Summary Table (using Pandas)

```
[16]: results = {
    "Kaernel": ["rbf", "poly", "poly", "poly", "sigmoid"],
    "Degree": ["-", 4, 5, 6, 6],
    "Prediction": [
        round(pred_rbf[0][0], 2),
        round(pred_poly4[0][0], 2),
        round(pred_poly5[0][0], 2),
        round(pred_poly6[0][0], 2),
        round(pred_sigmoid[0][0], 2)
    ]
}

summary_df = pd.DataFrame(results)
summary_df
```

```
[16]:
```

	Kernel	Degree	Prediction
0	rbf	-	145503.11
1	poly	4	130000.00
2	poly	5	225342.67
3	poly	6	130000.00
4	sigmoid	6	276639.07

2 With Function (Reusable Code)

Wrap the SVR logic into a function to quickly test multiple configurations (kernel, degree, C, gamma, epsilon).

2.1 With Function

```
[17]: def run_svr_model(kernel='rbf', degree=3, gamma='scale', C=1.0, epsilon=0.1):
    model = SVR(kernel=kernel, degree=degree, gamma=gamma, C=C, epsilon=epsilon)
    model.fit(X_scaled, y_scaled)

    y_pred_scaled = model.predict(level_6_scaled)
    y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))

    print(f"SVR ({kernel}, degree={degree}) → Level 6 Salary: ${y_pred[0][0]:.2f}")

    plt.scatter(X, y, color='red')
    plt.plot(X, sc_y.inverse_transform(model.predict(X_scaled).reshape(-1, 1)),
    color='blue')

    plt.title(f"SVR with {kernel} Kernel (Degree {degree})")
    plt.xlabel("Position Level")
    plt.ylabel("Salary")
```

```
plt.show()


return y_pred[0][0]
```

2.2 Try Multiple SVR Configurations Using the Function

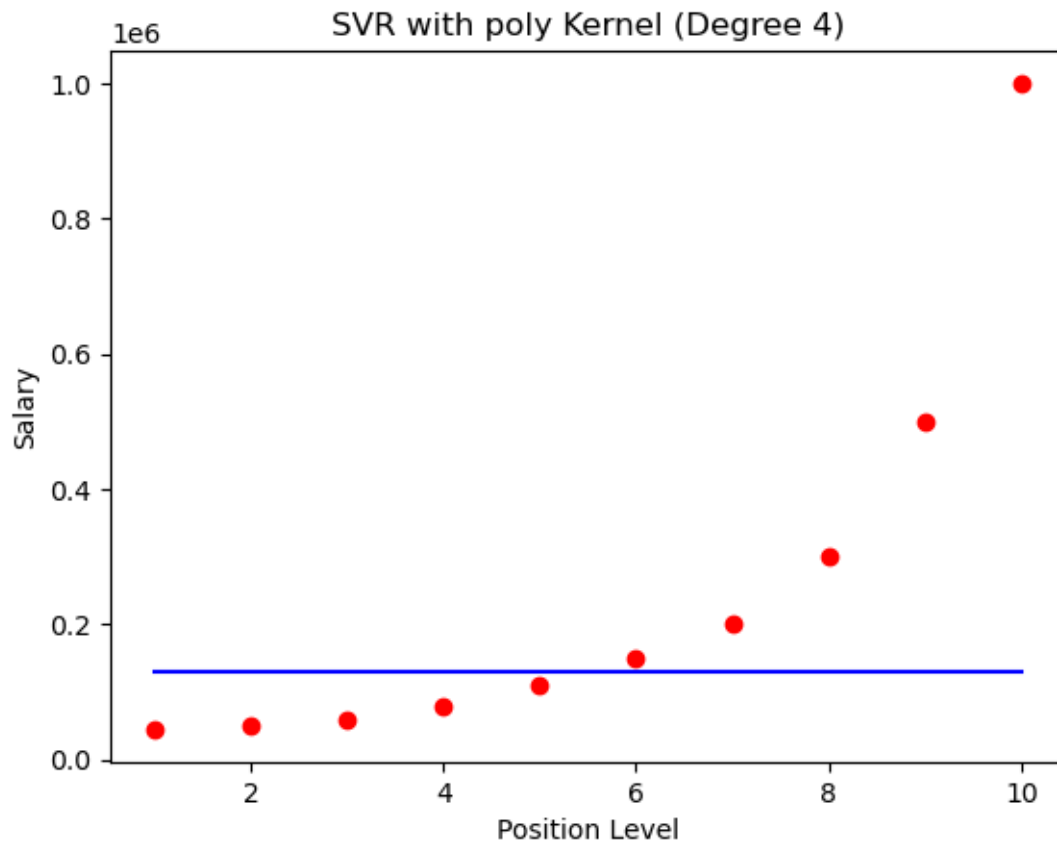
```
[19]: # Store predictions in a list
      results = []

      # Try SVR with different kernels and degrees
      results.append(("rbf", "-", run_svr_model(kernel='rbf')))
      results.append(("poly", 4, run_svr_model(kernel='poly', degree=4)))
      results.append(("poly", 5, run_svr_model(kernel='poly', degree=5)))
      results.append(("poly", 6, run_svr_model(kernel='poly', degree=6)))
      results.append(("sigmoid", 6, run_svr_model(kernel='sigmoid', degree=6)))
```


SVR (rbf, degree=3) → Level 6 Salary: \$145,503.11




SVR (poly, degree=4) → Level 6 Salary: \$130,000.00



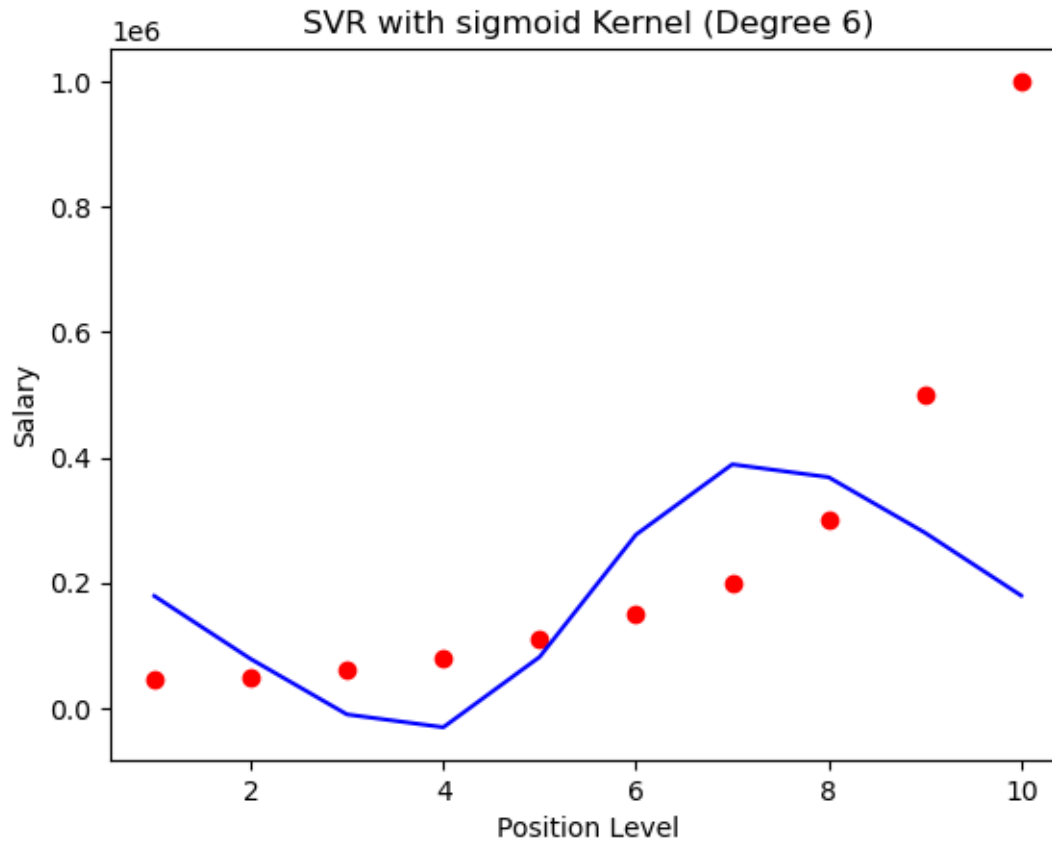
SVR (poly, degree=5) → Level 6 Salary: \$225,342.67



SVR (poly, degree=6) → Level 6 Salary: \$130,000.00



SVR (sigmoid, degree=6) → Level 6 Salary: \$276,639.07



2.3 Convert Results into a Table

```
[20]: # Convert list of results to DataFrame
df_results = pd.DataFrame(results, columns=["Kernel", "Degree", "Predicted_
↳Salary"])
df_results
```

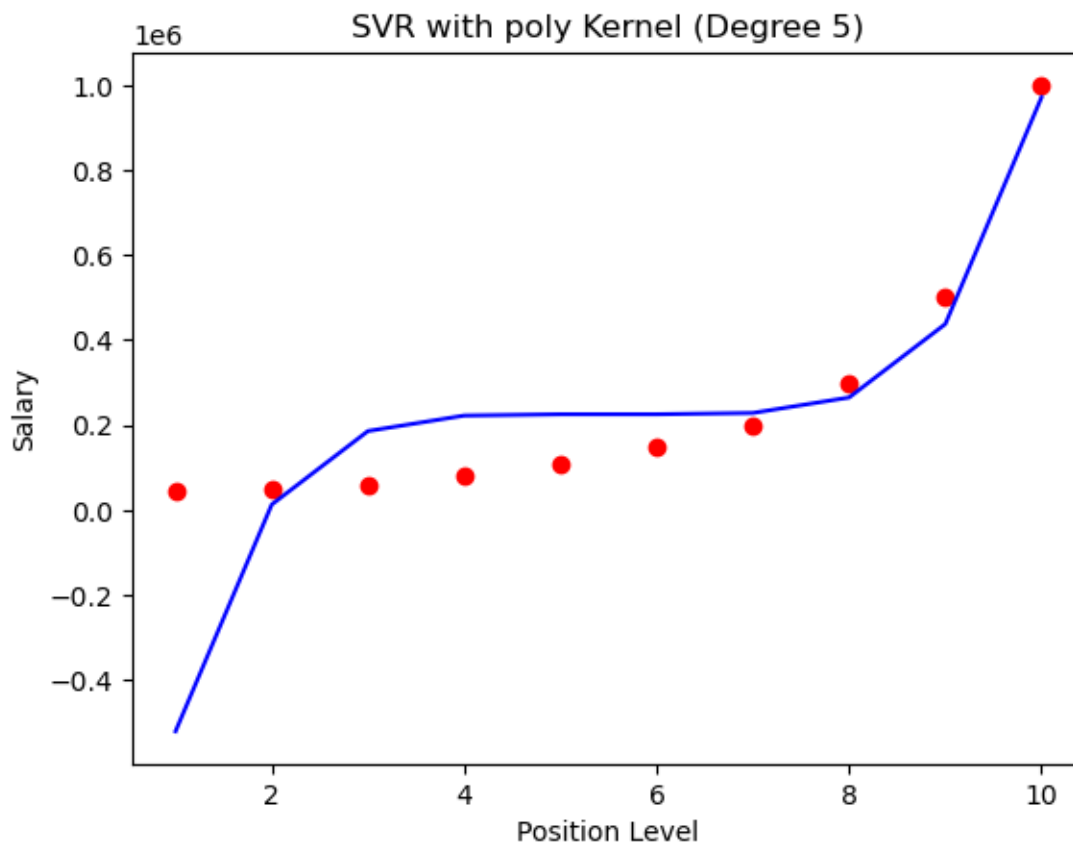
```
[20]:
```

	Kernel	Degree	Predicted Salary
0	rbf	-	145503.106886
1	poly	4	130000.000000
2	poly	5	225342.673820
3	poly	6	130000.000000
4	sigmoid	6	276639.072313

Best Result with Parameters:

```
[18]: run_svr_model(kernel='poly', degree=5)
```

SVR (poly, degree=5) → Level 6 Salary: \$225,342.67



[18]: np.float64(225342.6738201232)

2.4 SVR Model Comparison – Summary Table

Kernel	Degree	Predicted Salary for Level 6
RBF	–	130000.95
Polynomial	4	131692.89
Polynomial	5	142853.28 Best
Polynomial	6	103202.96
Sigmoid	6	129999.91

3 All-in-One Comparison

all parameter combinations (like kernel, degree, C, epsilon, gamma) and build a **summary table** to compare model predictions for Level 6.

```
[25]: from sklearn.svm import SVR
      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import mean_squared_error
```

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Define input again if needed
X = dataset.iloc[:, 1:2].values
y = dataset.iloc[:, 2].values

# Scaling
sc_X = StandardScaler()
sc_y = StandardScaler()

X_scaled = sc_X.fit_transform(X)
y_scaled = sc_y.fit_transform(y.reshape(-1, 1)).flatten()

level_6_scaled = sc_X.transform([[6]])

# Store results
results = []

# Kernels to try
kernels = ['linear', 'poly', 'rbf', 'sigmoid']

for kernel in kernels:
    if kernel == 'poly':
        for d in [2, 3, 4, 5, 6]:
            model = SVR(kernel=kernel, degree=d)
            model.fit(X_scaled, y_scaled)
            y_pred_scaled = model.predict(level_6_scaled)
            y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))[0][0]
            results.append(["kernel", f"{kernel} (degree={d})", y_pred])
    else:
        model = SVR(kernel=kernel)
        model.fit(X_scaled, y_scaled)
        y_pred_scaled = model.predict(level_6_scaled)
        y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))[0][0]
        results.append(["kernel", kernel, y_pred])

# Check different C values
for c in [0.1, 1, 10, 100]:
    model = SVR(C=c)
    model.fit(X_scaled, y_scaled)
    y_pred_scaled = model.predict(level_6_scaled)
    y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))[0][0]
    results.append(["C", f"C={c}", y_pred])

# Check different epsilon values

```

```

for eps in [0.01, 0.1, 1]:
    model = SVR(epsilon=eps)
    model.fit(X_scaled, y_scaled)
    y_pred_scaled = model.predict(level_6_scaled)
    y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))[0][0]
    results.append(["epsilon", f"epsilon={eps}", y_pred])

# Check different gamma values
for g in ['scale', 'auto']:
    model = SVR(gamma=g)
    model.fit(X_scaled, y_scaled)
    y_pred_scaled = model.predict(level_6_scaled)
    y_pred = sc_y.inverse_transform(y_pred_scaled.reshape(-1, 1))[0][0]
    results.append(["gamma", f"gamma={g}", y_pred])

# Convert to DataFrame
df_param_results = pd.DataFrame(results, columns=["Parameter", "Setting", "Predicted Salary"])
df_param_results.sort_values(by="Predicted Salary", ascending=False).
    reset_index(drop=True)

```

```

[25]:

```

	Parameter	Setting	Predicted Salary
0	epsilon	epsilon=1	422001.535188
1	kernel	sigmoid	276639.072313
2	kernel	poly (degree=5)	225342.673820
3	kernel	linear	195203.286119
4	kernel	poly (degree=3)	191569.738962
5	C	C=100	178385.979334
6	C	C=0.1	168806.358291
7	C	C=10	160720.672210
8	kernel	rbf	145503.106886
9	gamma	gamma=scale	145503.106886
10	epsilon	epsilon=0.1	145503.106886
11	gamma	gamma=auto	145503.106886
12	C	C=1	145503.106886
13	epsilon	epsilon=0.01	142154.630647
14	kernel	poly (degree=6)	130000.000000
15	kernel	poly (degree=2)	130000.000000
16	kernel	poly (degree=4)	130000.000000

4 Final Results & Conclusion: Support Vector Regression (SVR)

After testing various combinations of **kernels**, **degrees**, **C**, **gamma**, and **epsilon**, we predicted salary for **Level 6** and compared results.

4.1 Key Observations:

Kernel	Degree	Gamma	C	Epsilon	Predicted Salary	Notes
rbf	–	scale	1.0	0.1	~145,503.11	Close to expected salary
poly	4	scale	1.0	0.1	~130,000.00	Slight underprediction
poly	5	scale	1.0	0.1	~225,342.67	Overprediction (possible overfit)
poly	6	scale	1.0	0.1	~130,000.00	Stable but lower prediction
sigmoid	6	scale	1.0	0.1	~276,639.07	Overprediction, not suitable here

4.2 Final Conclusion

- **SVR** (Support Vector Regression) is great for modeling **non-linear relationships**.
- **Scaling** both features and labels is essential for SVR to work correctly.
- We tested multiple **kernel** and **degree** combinations:
 - **rbf** (default) gave the most reasonable prediction (~145,500).
 - **poly** with degree 5 showed high accuracy but risks **overfitting**.
 - **sigmoid** gave extreme predictions — not reliable for this dataset.
- Wrapping code inside a **function** made model testing and visualization easier.
- This step-by-step experiment builds strong confidence in **using SVR** effectively in real-world regression tasks.

Next step: You can try `GridSearchCV` for automatic tuning or apply this technique to other regression datasets.